

# Kernel Methods



Dept. Computer Science & Engineering,  
Shanghai Jiao Tong University

# Key Points in the Last Talk



- Good representation of function spaces
  - Easy to implement (efficient in space and time)
  - Good for generalization
  - Easy to select good models
- Good parameter for model selection
  - Effective degrees of freedom
  - CV for Model selection
- Reproducing Kernel Hilbert Space
  - Polynomial Kernel
- Spline & Wavelet



- One-Dimensional Kernel Smoothers
- Local Regression
- Local Likelihood
- Kernel Density estimation
- Naive Bayes
- Radial Basis Functions
- Mixture Models and EM

Course CS & 304H - STU Statistical Learning theory & Applications

# Objectives of OBE



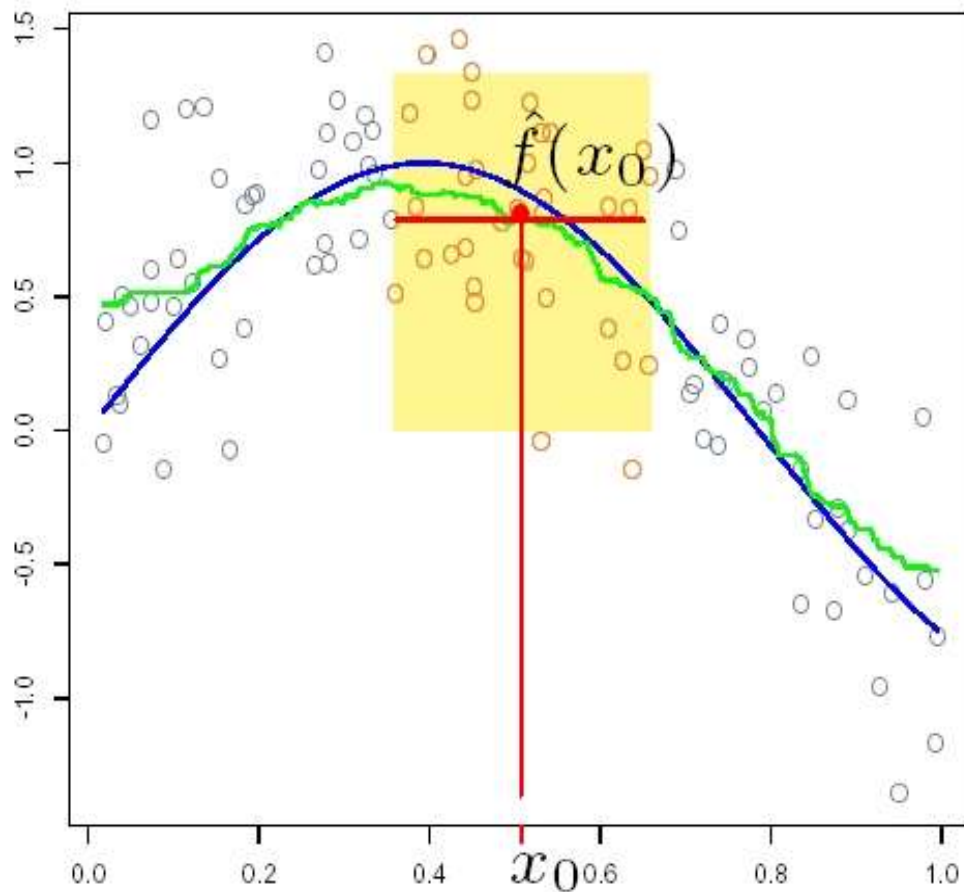
- To understand why increasing the order of polynomials causes the increase of model variance
- How to kernel method to implement local regression
- Probability density estimation by kernel methods
- EM algorithm for estimating GMM

Course CS &304H SJTU Statistical Learning Theory & Application

# One-Dimensional Kernel Smoothers



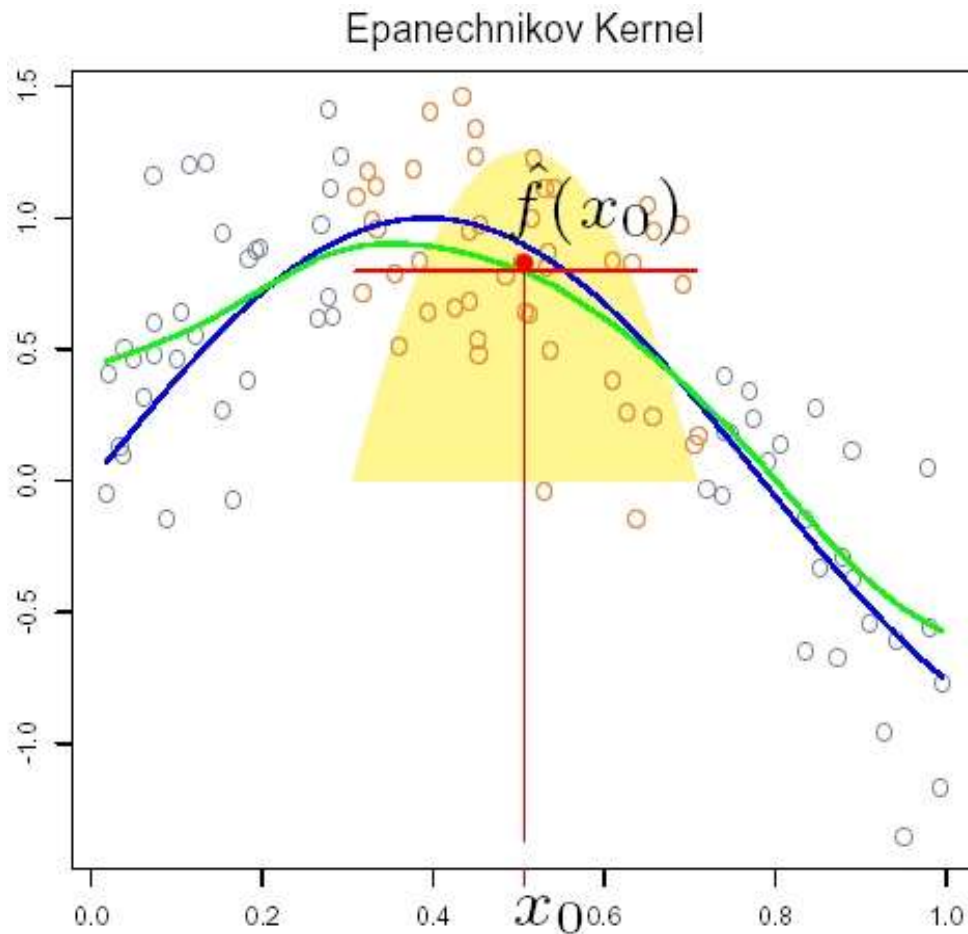
Nearest-Neighbor Kernel



- k-NN:  

$$\hat{f}(x) = Ave(y_i \mid x_i \in N_k(x))$$
- 30-NN curve is bumpy, since  $\hat{f}(x)$  is discontinuous in  $x$ .
- The average changes in a discrete way, leading to a discontinuous  $\hat{f}(x)$ .

# One-Dimensional Kernel Smoothers



- Nadaraya-Watson Kernel weighted average:

$$\hat{f}(x_0) = \frac{\sum_{i=1}^N K_{\lambda}(x_0, x_i) y_i}{\sum_{i=1}^N K_{\lambda}(x_0, x_i)}$$

- Epanechnikov quadratic kernel:

$$K_{\lambda}(x_0, x) = D\left(\frac{|x - x_0|}{\lambda}\right)$$

# One-Dimensional Kernel Smoothers



- More general kernel:

$$K_{\lambda}(x_0, x) = D \left( \frac{|x - x_0|}{h_{\lambda}(x_0)} \right)$$

- $h_{\lambda}(x_0)$  : width function that determines the width of the neighborhood at  $x_0$ .
- For quadratic kernel  $h_{\lambda}(x_0) = \lambda$ , Bias  $\approx$  constant
- For k-NN kernel  $\lambda \xleftrightarrow[\text{replaced}]{k} h_k(x_0) = |x_0 - x_{[k]}|$ ,  
Variance  $\approx$  constant

# One-Dimensional Kernel Smoothers



- Three popular kernel for local smoothing:

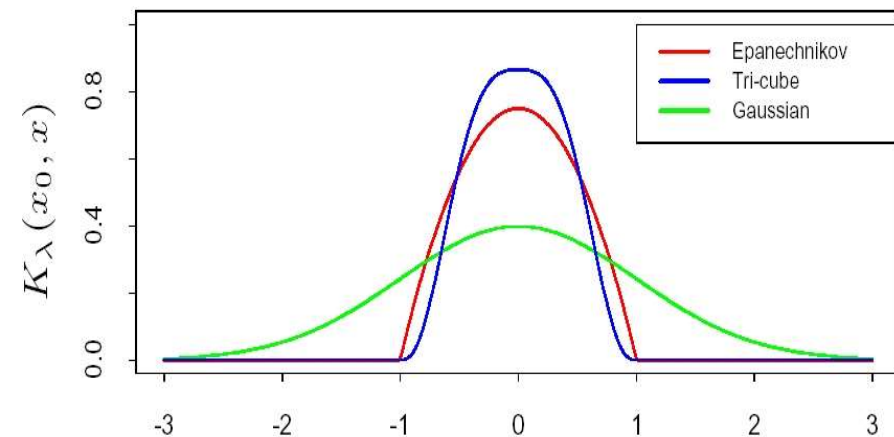
$$K_{\lambda}(x_0, x) = D\left(\frac{|x - x_0|}{\lambda}\right)$$

- Epanechnikov kernel and tri-cube kernel are compact but tri-cube has two continuous derivatives
- Gaussian kernel is infinite support

**Gaussian** :  $D(t) = \phi(t) = e^{-\frac{1}{2}t^2}$

**Epanechnikov** :  $D(t) = \begin{cases} \frac{3}{4} (1 - t^2) & \text{if } |t| < 1 \\ 0 & \text{otherwise} \end{cases}$

**Tri-Cube** :  $D(t) = \begin{cases} (1 - t^2)^3 & \text{if } |t| < 1 \\ 0 & \text{otherwise} \end{cases}$

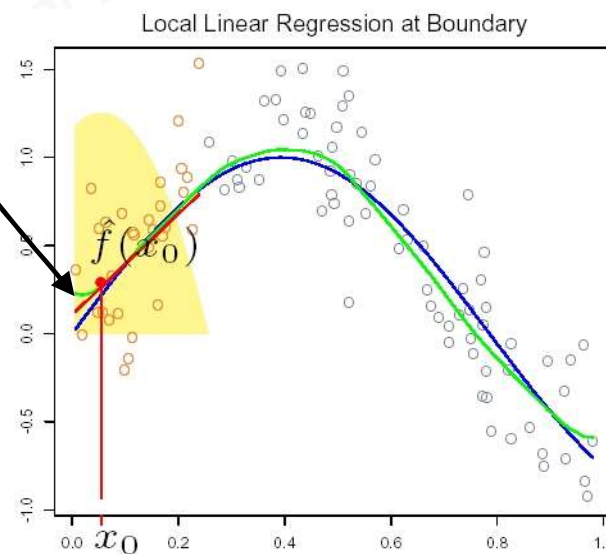
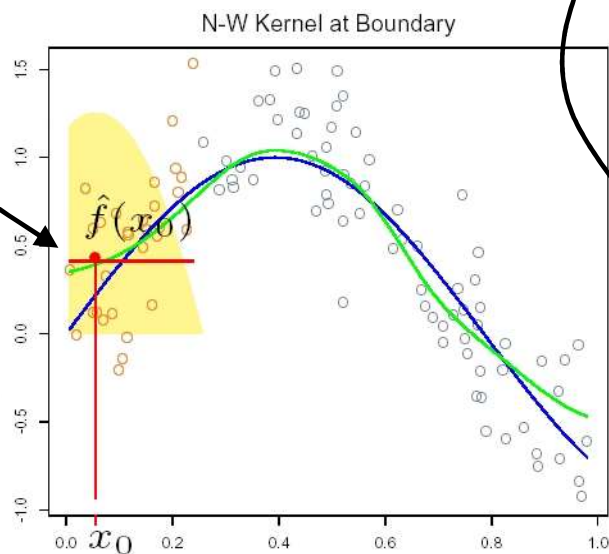




# Local Linear Regression



- Boundary issue
  - Badly biased on the boundaries because of the asymmetry of the kernel in the region.
  - Linear fitting remove the bias to first order



# Local Linear Regression



- Locally weighted linear regression make a first-order correction
- Separate weighted least squares at each target point  $x_0$ :

$$\min_{\alpha(x_0), \beta(x_0)} \sum_{i=1}^N K_{\lambda}(x_0, x_i) [y_i - \alpha(x_0) - \beta(x_0)x_i]^2$$

- The estimate:  $\hat{f}(x_0) = \hat{\alpha}(x_0) + \hat{\beta}(x_0)x_0$
- $b(x)^T = (1, x)$ ;  $B$ :  $N \times 2$  regression matrix with  $i$ -th row  $b(x_i)^T$ ;

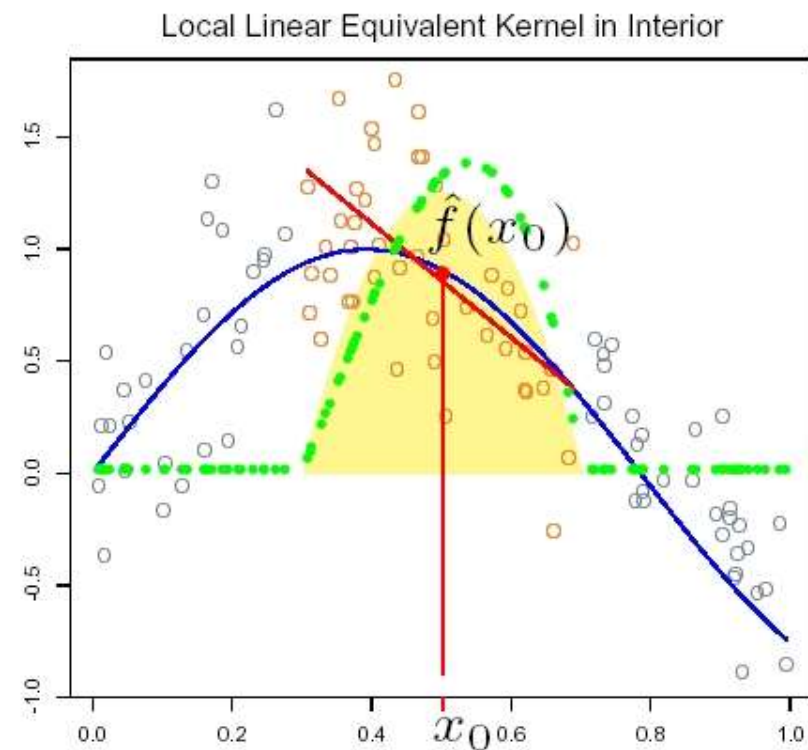
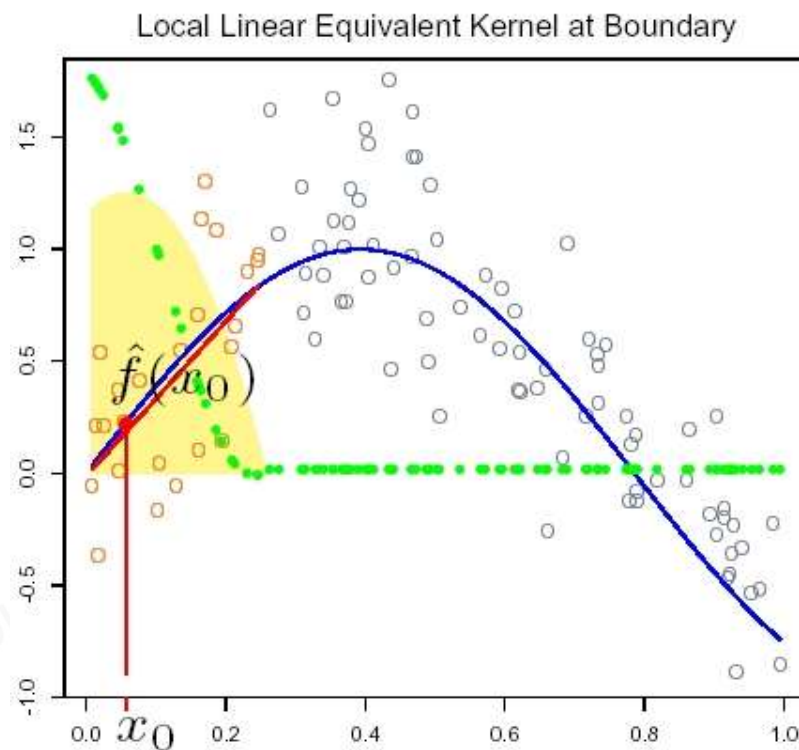
$$W_{N \times N}(x_0) = \text{diag}(K_{\lambda}(x_0, x_i)), i = 1, \dots, N$$

$$\hat{f}(x_0) = b(x_0)^T (B^T W(x_0) B)^{-1} B^T W(x_0) y = \sum_{i=1}^N l_i(x_0) y_i$$

# Local Linear Regression



- The weights  $l_i(x_0)$  combine the weighting kernel  $K_\lambda(x_0, \cdot)$  and the least squares operations——**Equivalent Kernel**



# Local Linear Regression



- The expansion for  $E\hat{f}(x_0)$ , using the linearity of local regression and a series expansion of the true function  $f$  around  $x_0$

$$E\hat{f}(x_0) = \sum_{i=1}^N l_i(x_0) f(x_i) = f(x_0) \sum_{i=1}^N l_i(x_0) + f'(x_0) \sum_{i=1}^N (x_i - x_0) l_i(x_0) + \frac{f''(x_0)}{2} \sum_{i=1}^N (x_i - x_0)^2 l_i(x_0) + R$$

$$\sum_{i=1}^N l_i(x_0) = 1, \quad \sum_{i=1}^N (x_i - x_0) l_i(x_0) = 0$$

## For local regression

- The bias  $E\hat{f}(x_0) - f(x_0)$  depends only on quadratic and higher-order terms in the expansion of  $f$ .

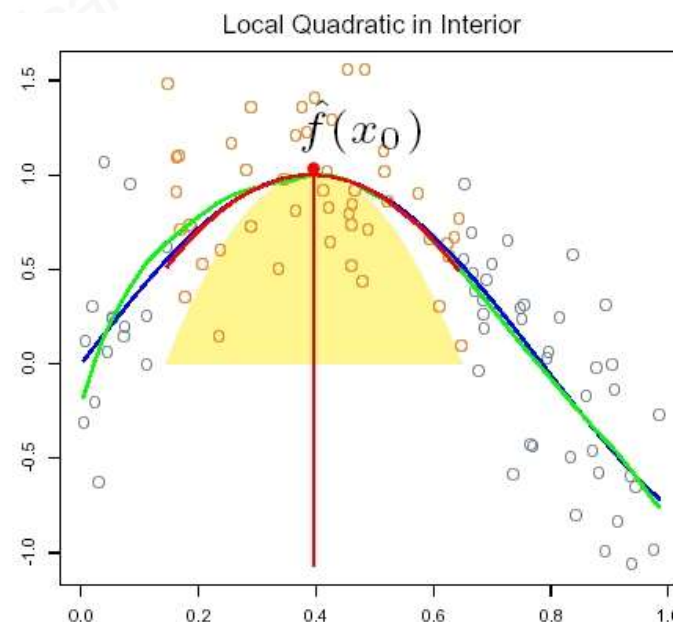
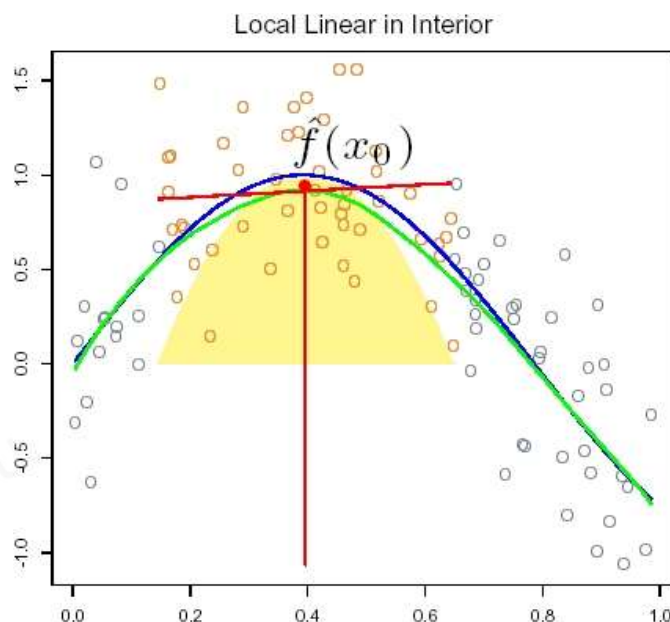
# Local Polynomial Regression



- Fit local polynomial fits of any degree  $d$

$$\min_{\alpha(x_0), \beta_j(x_0), j=1, \dots, d} \sum_{i=1}^N K_{\lambda}(x_0, x_i) \left[ y_i - \alpha(x_0) - \sum_{j=1}^d \beta_j(x_0) x_i^j \right]^2$$

$$\hat{f}(x_0) = \hat{\alpha}(x_0) + \sum_{j=1}^d \hat{\beta}_j(x_0) x_0^j$$

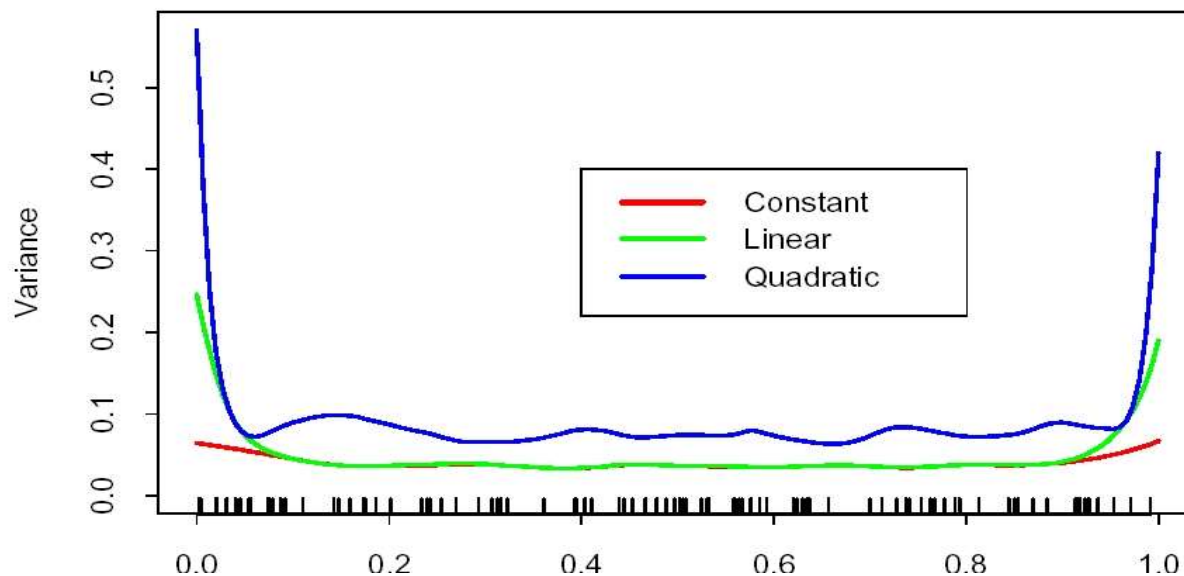


# Local Polynomial Regression



- Bias only have components of degree  $d+1$  and higher.
- The reduction for bias costs the increased variance.

$$\text{var}(\hat{f}(x_0)) = \sigma^2 \|l(x_0)\|^2, \quad \|l(x_0)\| \text{ increases with } d$$



# Kernel Width Selection



- $\lambda$  is a parameter in kernel  $K_\lambda$ , which controls the width of the kernel.
  - $\lambda$  takes the **radius** of supporting region for compact supporting kernel;
  - $\lambda$  takes the **variance** for Gaussian Kernel;
  - $\lambda$  takes **k/N** for KNN method.
- Kernel width is related to model selection
  - Wide width leads to large bias and small var.
  - Narrow width leads to small bias and large var.

# Structured Local Regression



- Structured kernels

$$K_{\lambda, A}(x_0, x) = D \left( \frac{(x - x_0)^T A (x - x_0)}{\lambda} \right)$$

- Introduce structure by imposing appropriate restrictions on A

- Structured regression function

$$f(X_1, X_2, \dots, X_p) = \alpha + \sum_j g_j(X_j) + \sum_{k < l} g_{kl}(X_k, X_l) + \dots$$

- Introduce structure by eliminating some of the higher-order terms



- Any parametric model can be made local:

- Parameter associated with  $y_i$  :  $\theta_i = \theta(x_i) = x_i^T \beta$

- Log-likelihood:  $l(\beta) = \sum_{i=1}^N l(y_i, x_i^T \beta)$

- Model  $\theta(X)$  likelihood local to  $x_0$  :

$$l(\beta(x_0)) = \sum_{i=1}^N K_{\lambda}(x_0, x_i) l(y_i, x_i^T \beta(x_0))$$

- A varying coefficient model  $\theta(z)$

$$l(\theta(z_0)) = \sum_{i=1}^N K_{\lambda}(z_0, z_i) l(y_i, \eta(x_i, \theta(z_0)))$$

$$e.g. \quad \eta(x, \theta) = x^T \theta$$

- Logistic Regression

$$\Pr(G = j \mid X = x) = \frac{\exp(\beta_{j0} + \beta_j^T x)}{1 + \sum_{k=1}^{J-1} \exp(\beta_{k0} + \beta_k^T x)}$$

- Local log-likelihood for the J class model

$$\sum_{i=1}^N K_{\lambda}(x_0, x_i) \left\{ \beta_{g_i 0}(x_0) + \beta_{g_i}(x_0)^T (x_i - x_0) - \log \left[ 1 + \sum_{k=1}^{J-1} \exp(\beta_{k0}(x_0) + \beta_k(x_0)^T (x_i - x_0)) \right] \right\}$$

- Center the local regressions at

$$\hat{\Pr}(G = j \mid X = x) = \frac{\exp(\hat{\beta}_{j0}(x_0))}{1 + \sum_{k=1}^{J-1} \exp(\hat{\beta}_{k0}(x_0))}$$

# Kernel Density Estimation



- A natural local estimate

$$\hat{f}_X(x_0) = \frac{\#x_i \in N(x_0)}{N\lambda}$$

- The smooth Parzen estimate

$$\hat{f}_X(x_0) = \frac{1}{N\lambda} \sum_{i=1}^N K_\lambda(x_0, x_i)$$

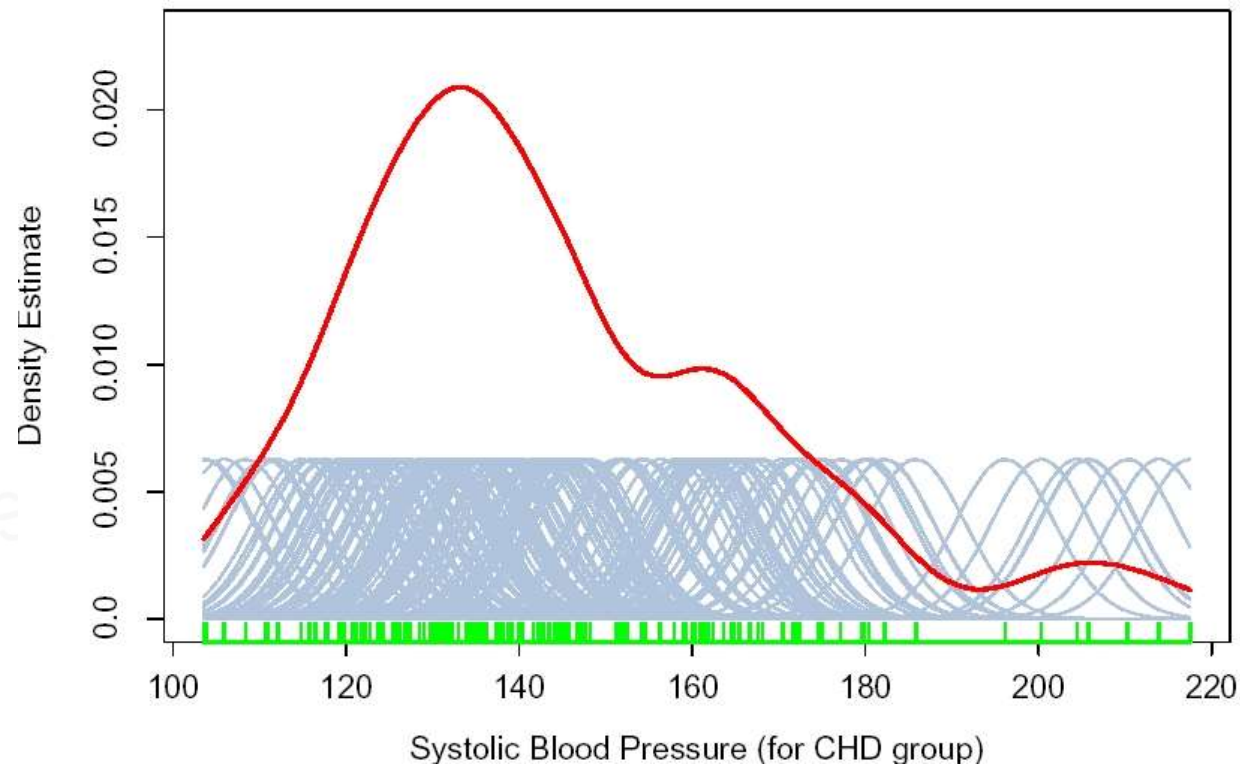
- For Gaussian kernel  $K_\lambda(x_0, x_i) / \lambda = \phi(x_i - x_0)$
- The estimate become

$$\begin{aligned} \hat{f}_X(x) &= \frac{1}{N} \sum_{i=1}^N \phi_\lambda(x_i - x_0) \\ &= \frac{1}{N(2\lambda^2\pi)^{p/2}} \sum_{i=1}^N \exp\left(-\frac{1}{2}(\|x_i - x_0\|/\lambda)^2\right) \end{aligned}$$

# Kernel Density Estimation



- A kernel density estimate for systolic blood pressure. The density estimate at each point is the average contribution from each of the kernels at that point.



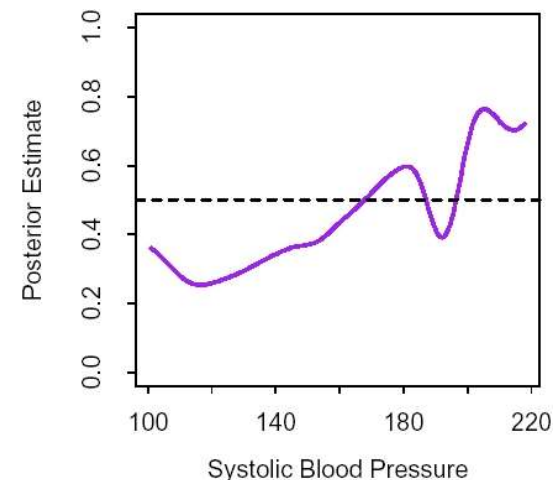
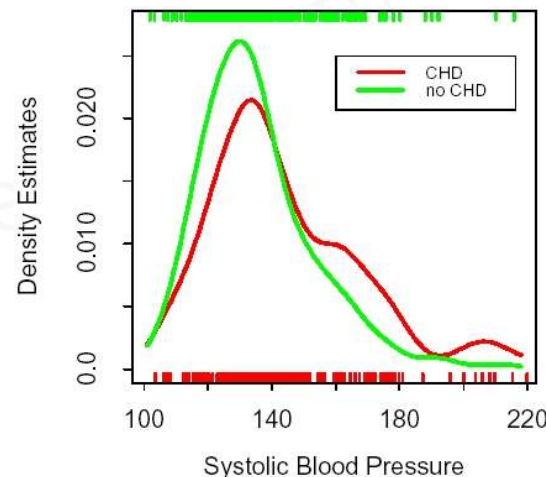
# Kernel Density Classification



- Bayes' theorem: 
$$\hat{\Pr}(G = j | X = x_0) = \frac{\hat{\pi}_j \hat{f}_j(x_0)}{\sum_{k=1}^J \hat{\pi}_k \hat{f}_k(x_0)}$$

$$f_j(x) = \Pr(X = x | G = j)$$

- The estimate for CHD (Coronary heart disease 冠心病) uses a *Gaussian kernel density estimate*.



# Kernel Density Classification



- The population class densities and the posterior probabilities

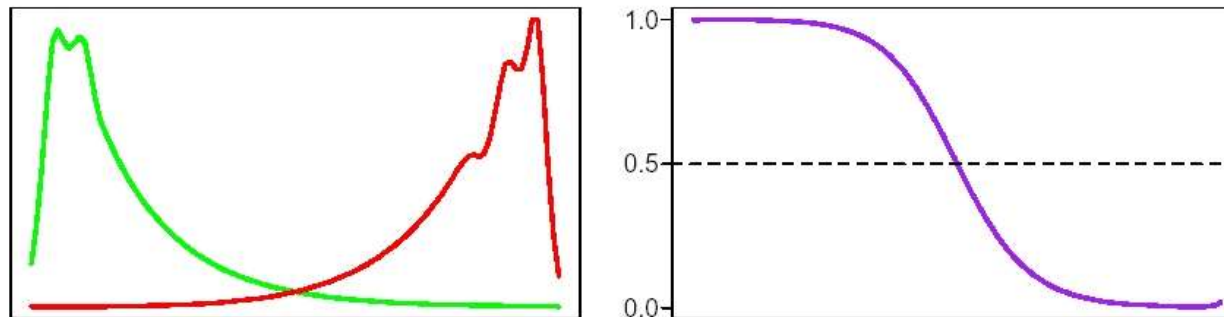


Figure 6.15: *The population class densities may have interesting structure (left) that disappears when the posterior probabilities are formed (right).*

# Naïve Bayes



- Naïve Bayes model assume that given a class  $G=j$ , the features  $X_k$  are independent:

$$f_j(X) = \prod_{k=1}^p f_{jk}(X_k)$$

- $\hat{f}_{jk}(X_k)$  is kernel density estimate, or Gaussian, for coordinate  $X_k$  in class  $j$ .
- If  $X_k$  is categorical, use Histogram.

$$\begin{aligned} \text{logit} \frac{\Pr(G = \ell | X)}{\Pr(G = J | X)} &= \log \frac{\pi_\ell f_\ell(X)}{\pi_J f_J(X)} = \log \frac{\pi_\ell \prod_{k=1}^p f_{\ell k}(X_k)}{\pi_J \prod_{k=1}^p f_{Jk}(X_k)} \\ &= \log \frac{\pi_\ell}{\pi_J} + \sum_{k=1}^p \log \frac{f_{\ell k}(X_k)}{f_{Jk}(X_k)} = \alpha_\ell + \sum_{k=1}^p g_{\ell k}(X_k) \end{aligned}$$

# Radial Basis Function & Kernel



- Radial basis function combine the local and flexibility of kernel methods.

$$f(x) = \sum_{j=1}^M K_{\lambda_j}(\xi_j, x) \beta_j = \sum_{j=1}^M D\left(\frac{\|x - \xi_j\|}{\lambda_j}\right) \beta_j$$

- Each basis element is indexed by a location or prototype parameter  $\xi_j$  and a scale parameter  $\lambda_j$
- $D$  , a pop choice is the standard Gaussian density function.



# Radial Basis Function & Kernel



- For simplicity, focus on least squares methods for regression, and use the Gaussian kernel.
- RBF network model:

$$\min_{\{\lambda_j, \xi_j, \beta_j\}_1^M} \sum_{i=1}^N \left( y_i - \beta_0 - \sum_{j=1}^M \beta_j \exp \left\{ -\frac{(x_i - \xi_j)^T (x_i - \xi_j)}{\lambda_j^2} \right\} \right)^2$$

- Estimate the  $\{\lambda_j, \xi_j\}$  separately from the  $\beta_j$ .
- A undesirable side effect of creating holes——regions of  $\mathbb{R}^p$  where none of the kernels has appreciable support.

# Radial Basis Function & Kernel

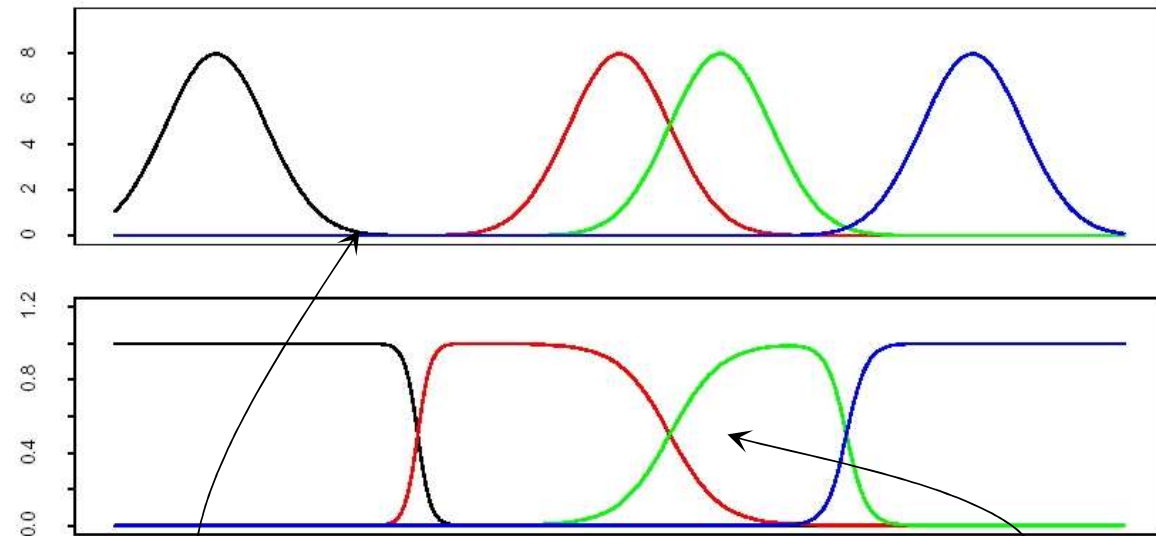


- Renormalized radial basis functions.

$$h_j(x) = \frac{D(\|x - \xi_j\| / \lambda)}{\sum_{k=1}^M D(\|x - \xi_k\| / \lambda)}$$

- The expansion in renormalized

$$\begin{aligned} f(x) &= \sum_{i=1}^N y_i \frac{K(x_0, x_i)}{\sum_{i=1}^N K_\lambda(x_0, x_i)} \\ &= \sum_{i=1}^N y_i h_i(x_0) \end{aligned}$$



Gaussian radial basis function with fixed width can leave holes. Renormalized Gaussian radial basis function produce basis functions similar in some respects to B-splines.

# Mixture Models & EM



- Gaussian Mixture Model

$$f(x) = \sum_{m=1}^M \alpha_m \phi(x; \mu_m, \Sigma_m)$$

- $\alpha_m$  are mixture proportions,  $\sum_{m=1}^M \alpha_m = 1$

- EM algorithm for mixtures

- Given  $x_1, x_2, \dots, x_n$ , log-likelihood:

$$l(y, \theta) = \sum_{i=1}^N \log[\alpha \phi_{\theta_1}(x_i) + (1 - \alpha) \phi_{\theta_2}(x_i)]$$

Bad

- Suppose we observe Latent Binary

$$L(x, z, \theta) = \sum_{\substack{i=1 \\ z_i=1}}^N \log[\alpha \phi_{\theta_1}(x_i)] + \sum_{\substack{i=1 \\ z_i=0}}^N \log[(1 - \alpha) \phi_{\theta_2}(x_i)]$$

$z$  such that  $z = 1 \Rightarrow x \sim \phi_{\theta_1}$ ,  $z = 0 \Rightarrow x \sim \phi_{\theta_2}$

Good

- Given  $\theta^0$ , compute

$$\theta^* = \arg \max \tilde{\ell}(\theta) = E[\ell(x, z, \theta) | (\theta^0, y)]$$

- For each sample

$$E(z_i | x_i, \theta^0) = \frac{\hat{\alpha} \phi_{\hat{\theta}_1}(x_i)}{\hat{\alpha} \phi_{\hat{\theta}_1}(x_i) + (1 - \hat{\alpha}) \phi_{\hat{\theta}_2}(x_i)} = w_i$$

$$\ell(\theta) = \sum_{i=1}^N w_i \log[\hat{\alpha} \phi_{\theta_1}(x_i)] + (1 - w_i) \log[(1 - \alpha) \phi_{\theta_2}(x_i)]$$

# The EM Algorithm



- The EM algorithm for two-component Gaussian mixtures
  - Take initial guesses  $\hat{\pi}, \hat{\mu}_1, \hat{\theta}_1, \hat{\mu}_2, \hat{\theta}_2$  for the parameters
  - Expectation Step: Compute the responsibilities

$$\hat{\gamma}_i = \frac{\hat{\pi} \phi_{\hat{\theta}_2}(y_i)}{(1 - \hat{\pi}) \phi_{\hat{\theta}_1}(y_i) + \hat{\pi} \phi_{\hat{\theta}_2}(y_i)}, \quad i = 1, \dots, N$$

$$\bullet \quad E(z_i | x_i, \theta^0) = \frac{\hat{\alpha} \phi_{\hat{\theta}_1}(x_i)}{\hat{\alpha} \phi_{\hat{\theta}_1}(x_i) + (1 - \hat{\alpha}) \phi_{\hat{\theta}_2}(x_i)} = w_i$$

$$\ell(\theta) = \sum_{i=1}^N w_i \log [\hat{\alpha} \phi_{\theta_1}(x_i)] + (1 - w_i) \log [(1 - \alpha) \phi_{\theta_2}(x_i)]$$

# The EM Algorithm

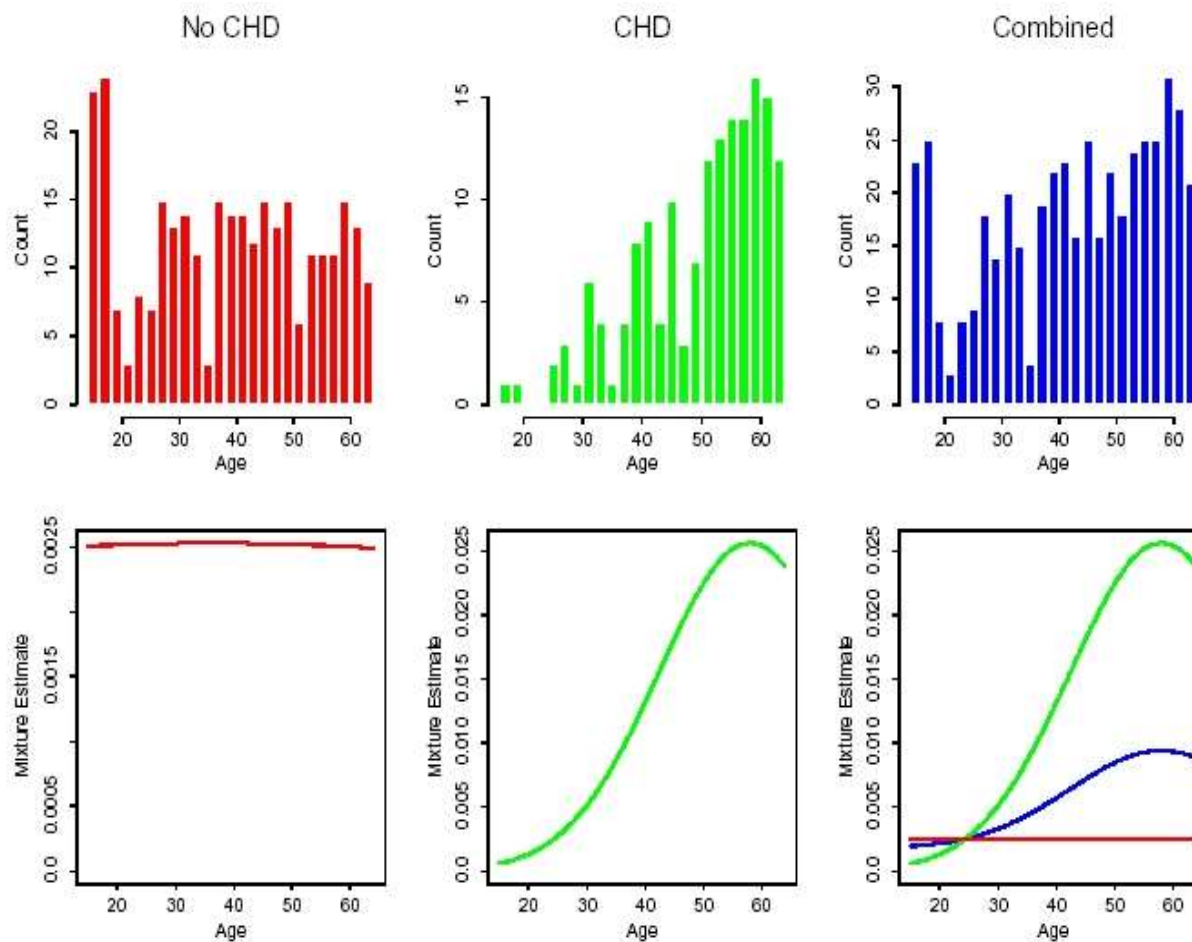


- Maximization Step: Compute the weighted means and variances

$$\begin{aligned}\hat{\mu}_1 &= \frac{\sum_{i=1}^N (1 - \hat{\gamma}_i) y_i}{\sum_{i=1}^N (1 - \hat{\gamma}_i)}, & \hat{\sigma}_1^2 &= \frac{\sum_{i=1}^N (1 - \hat{\gamma}_i) (y_i - \hat{\mu}_1)^2}{\sum_{i=1}^N (1 - \hat{\gamma}_i)}, \\ \hat{\mu}_2 &= \frac{\sum_{i=1}^N \hat{\gamma}_i y_i}{\sum_{i=1}^N \hat{\gamma}_i}, & \hat{\sigma}_2^2 &= \frac{\sum_{i=1}^N \hat{\gamma}_i (y_i - \hat{\mu}_2)^2}{\sum_{i=1}^N \hat{\gamma}_i}, \\ \hat{\pi} &= \sum_{i=1}^N \hat{\gamma}_i / N\end{aligned}$$

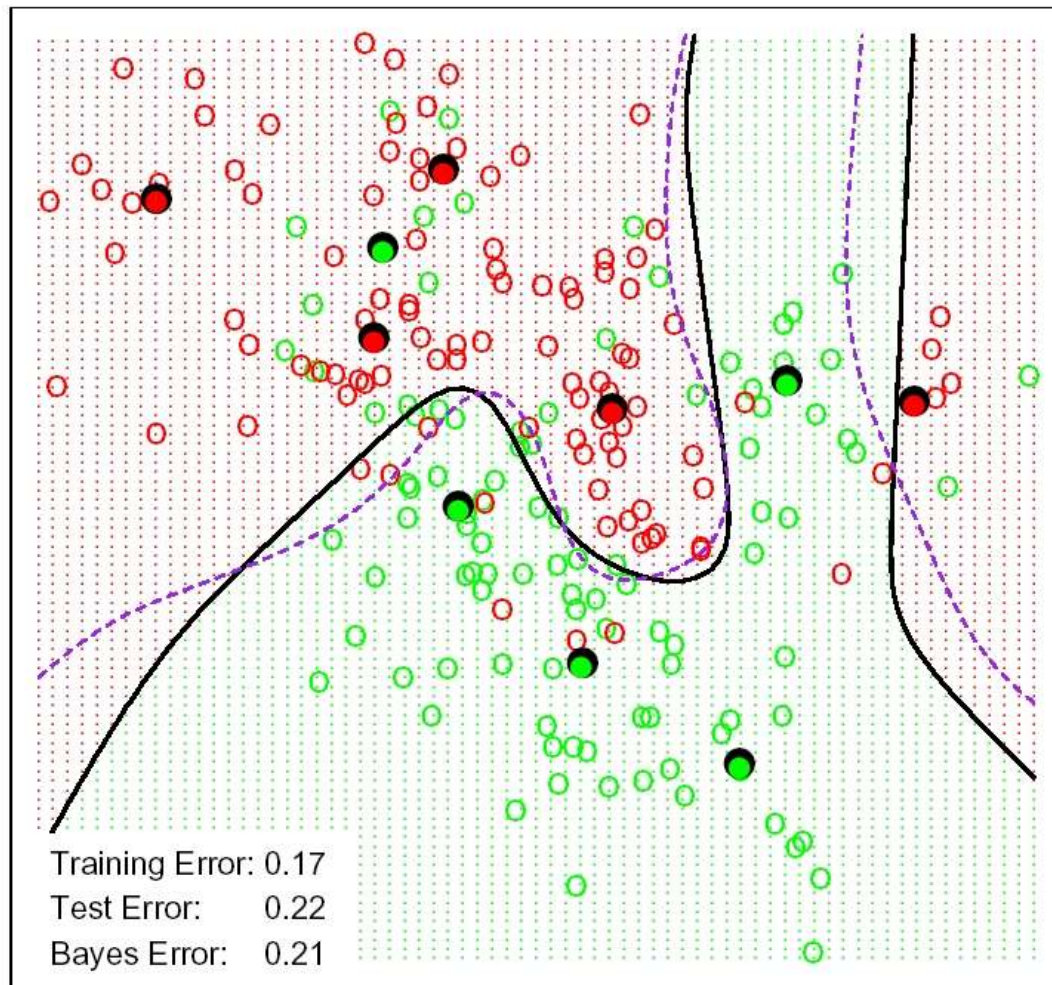
- Iterate 2 and 3 until convergence

# Mixture Models & EM



- Application of mixtures to the heart disease risk factor study.

# Mixture Models & EM



- Mixture model used for classification of the simulated data



# Summary



- To understand why increasing the order of polynomials causes the increase of model variance
- How to kernel method to implement local regression
- Probability density estimation by kernel methods
- EM algorithm for estimating GMM

**The End of Talk!**

---