



Talk 12 Unsupervised Learning

Part 2

Unsupervised Learning




- 1. Introduction**
- 2. Association Rules & Cluster Analysis**
-
- 4. Self-Organizing Maps**
- 5. Principal Components, Curves and Surfaces**
- 6. Non-negative Matrix Factorization**
- 7. Independent Component Analysis**
- 8. Multidimensional Scaling**
- 9. Nonlinear Dimension Reduction**
- 10. The Google PageRank Algorithm**

6. Independent Component Analysis (ICA)



- Speech Separation (Cocktail Party Problem)




Mathematical Formulation



Mixing Model

$$x_i(k) = a_{i1}s_1(k) + a_{i2}s_2(k) + \cdots + a_{in}s_n(k) + \varepsilon(k)$$

$$\mathbf{x}(k) = \mathbf{A}\mathbf{s}(k) + \boldsymbol{\varepsilon}(k)$$



- $\mathbf{s}(k) = (s_1(k), \dots, s_n(k))^T$: the vector of n-source signals;
- $\mathbf{x}(k) = (x_1(k), \dots, x_m(k))^T$: the vector of m-sensor signals;
- $\boldsymbol{\varepsilon}(k)$: the vector of sensor noises.
- \mathbf{A} is the mixing matrix.


Demixing Model



Problem: to estimate the source signals (or event-related potentials)
by using the sensor signals

$$\mathbf{y}(k) = \mathbf{W} \mathbf{x}(k)$$

- $\mathbf{y}(k) = (y_1(k), \dots, y_m(k))^T$: the vector of recovered signals
- \mathbf{W} is the demixing matrix.



- **Assumption:**

- ✓ The source signals are mutually independent

- **Model:**

- ✓ Linear instantaneous mixture
- ✓ Linear convolutive mixture

- **Principles**

- ✓ Maximum Entropy
- ✓ Minimum Mutual Information
- ✓ Joint Diagonalization of Cross-correlations
- ✓ Linear Predictability
- ✓ Sparseness Maximization

6. Independent Component Analysis (ICA)



- One of Approaches to recover independent component one by one:
 - Maximize the **NonGaussianity**
- Let mixing model: $\mathbf{x} = \mathbf{A}\mathbf{s} + \boldsymbol{\varepsilon}; \quad \mathbf{s} = (s_1, s_2, \dots, s_n)^T$
- We attempt to extract a single component from \mathbf{x}

$$y(k) = \mathbf{b}^T \mathbf{x}(k) = \mathbf{b}^T \mathbf{A}\mathbf{s}(k)$$

- Let $\mathbf{z} = \mathbf{A}^T \mathbf{b}$

$$y(k) = \mathbf{b}^T \mathbf{A}\mathbf{s}(k) = \mathbf{z}^T \mathbf{s}(k)$$

What we can derive from this observation?
CLT provides us a criterion for the ICA.

CLT – Central Limit theory



Assume X_1, X_2, \dots, X_n are iid samples from a probability with mean μ and variance σ^2 , define

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i; \quad U_n = \frac{\bar{X} - \mu}{\sigma / \sqrt{n}}$$

Then the random variable U_n converges to the normal distribution as the sample size n tends to infinite.

$$U_n = \frac{\bar{X} - \mu}{\sigma / \sqrt{n}} \rightarrow N(0, 1)$$

6. Independent Component Analysis (ICA)




- **Kurtosis is a measure of non-gaussianity**

$$\text{kurtosis}(y) = E[y^4] - 3(E[y^2])^2$$

$$\hat{b} = \text{var} \max_b |\text{kurtosis}(b^T \mathbf{x})|$$

FastICA algorithm is derived



$$\text{kurtosis}(y) = E[y^4] - 3(E[y^2])^2 \begin{cases} > 0, & \text{for SuperGaussian} \\ = 0, & \text{for Gaussian} \\ < 0, & \text{for SubGaussian} \end{cases}$$

Cost Function




Kullback-Leibler (KL) divergence between

$$p(y_1, \dots, y_K) \text{ and } \prod_{k=1}^K q_k(y_k)$$

$$KL(W) = \int p(\mathbf{y}) \log \frac{p(\mathbf{y})}{\prod_{k=1}^K q_k(y_k)} d\mathbf{y}$$

$$= \underbrace{-H(\mathbf{Y}; W)}_{\text{1. Joint Entropy of } \mathbf{y}} + \underbrace{\sum_{k=1}^K H(Y_k; W)}_{\text{2. Sum of marginal entropy of } y_k}$$



• Minimized when y_k are mutually independent

Cost Function



Assume that $\mathbf{y} = \mathbf{W}\mathbf{x}$, the differential entropy can be expressed by

$$H(\mathbf{y}) = H(\mathbf{x}) + \log |\det(\mathbf{W})|$$

Assume that $\mathbf{y} = \mathbf{W}\mathbf{x}$, \mathbf{W} is an invertable matrix, the probability density function of \mathbf{y} is given by $p_y(\mathbf{y}) = |\det(\mathbf{W})|^{-1} p_x(\mathbf{x})$

$$\begin{aligned} H(\mathbf{y}) &= -\int p(\mathbf{y}) \log(p(\mathbf{y})) d\mathbf{y} \\ &= -\int |\det(\mathbf{W})|^{-1} p_x(\mathbf{x}) \log(|\det(\mathbf{W})|^{-1} p_x(\mathbf{x})) |\det(\mathbf{W})| d\mathbf{x} \\ &= -\int p_x(\mathbf{x}) \log(p_x(\mathbf{x})) d\mathbf{x} + \log(|\det(\mathbf{W})|) \end{aligned}$$

Cost Function



Assume that $\mathbf{y} = \mathbf{W}\mathbf{x}$, the differential entropy can be expressed by

$$H(\mathbf{y}) = H(\mathbf{x}) + \log |\det(\mathbf{W})|$$

$$KL(W) = \int p(\mathbf{y}) \log \frac{p(\mathbf{y})}{\prod_{k=1}^K p(y_k)} d\mathbf{y}$$

$$= -H(\mathbf{Y}; W) + \sum_{k=1}^K H(Y_k; W)$$

$$= -H(x) - \log |\det(W)| - \sum_{k=1}^K E[\log(p_k(y_k))]$$

Gradient Descent



To update \mathbf{W} along the negative gradient of $KL(\mathbf{W})$

$$\begin{aligned}\Delta \mathbf{W} &\propto -\frac{\partial KL(\mathbf{W})}{\partial \mathbf{W}} = \left((\mathbf{W}^T)^{-1} - \int p(x) \phi(\mathbf{y}) \mathbf{x}^T dx \right) \\ &= \left((\mathbf{W}^T)^{-1} - \mathbb{E}_x [\phi(\mathbf{y}) \mathbf{x}^T] \right) \\ &= \left(\mathbf{I} - \mathbb{E}_y [\phi(\mathbf{y}) \mathbf{y}^T] \right) (\mathbf{W}^T)^{-1}\end{aligned}$$

.....

Nonlinear Function 2 \Rightarrow To be diagonalized

where

$$\phi(\mathbf{y}) \equiv \left[\frac{\partial \log p(y_1)}{\partial y_1}, \dots, \frac{\partial \log p(y_K)}{\partial y_K} \right]^T$$

This can be approximated by Sigmoid Function in speech signal.

Gradient Descent



To update \mathbf{W} along the negative gradient of $KL(W)$

$$\Delta \mathbf{W} = \alpha \left(\mathbf{I} - \mathbb{E}_y \left[\phi(\mathbf{y}) \mathbf{y}^T \right] \right) (\mathbf{W}^T)^{-1}$$

Modified Learning Algorithm

$$\phi(\mathbf{y}) \equiv \left[\frac{\partial \log p(y_1)}{\partial y_1}, \dots, \frac{\partial \log p(y_K)}{\partial y_K} \right]^T$$

$$\Delta \mathbf{W} = \alpha \left(\mathbf{I} - \mathbb{E}_y \left[\phi(\mathbf{y}) \mathbf{y}^T \right] \right) \mathbf{W}$$

The natural gradient of Nonsingular Matrix Manifold

$$\nabla l(\mathbf{W}) = \frac{\partial l(\mathbf{W})}{\partial \mathbf{W}} \mathbf{W}^T \mathbf{W}$$

ICA theoretical problems



- Convergence of gradient descent
 - Depends on the choice of activation function
 - For super Gaussian, $\varphi(y) = \tanh(y)$
 - For sub Gaussian, $\varphi(y) = y^3$
- Adaptation of activation function
 - Using generalized Gaussian family


$$\Delta \mathbf{W} = \alpha \left(\mathbf{I} - \mathbf{E}_y \left[\varphi(\mathbf{y}) \mathbf{y}^T \right] \right) \mathbf{W}$$

$$f(x) = \frac{c_p}{\alpha} \exp \left(-\frac{|x - \mu|^p}{2\alpha^p} \right), \quad c_p = \frac{p}{2^{(p+1)/p} \Gamma(1/p)}, \quad p > 0$$


- Determination of the number of ICs

ICA Solution (Non-Uniqueness)


Source S




Data X



PCA Solution



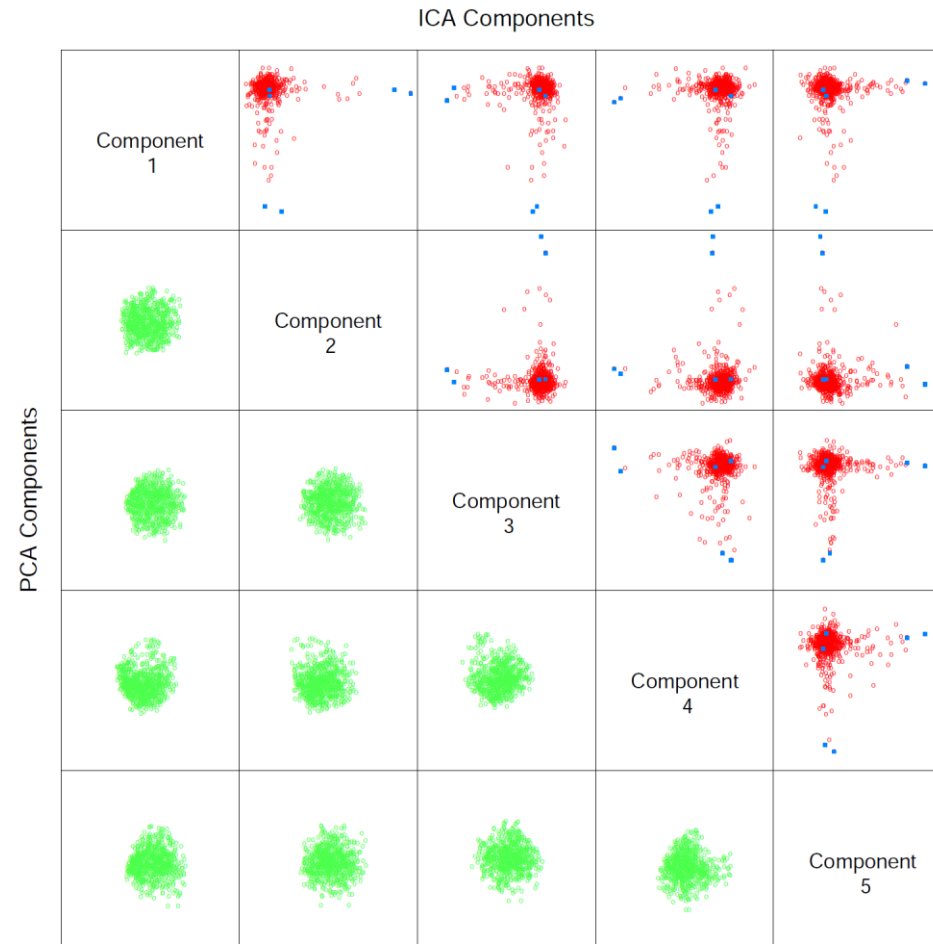
ICA Solution




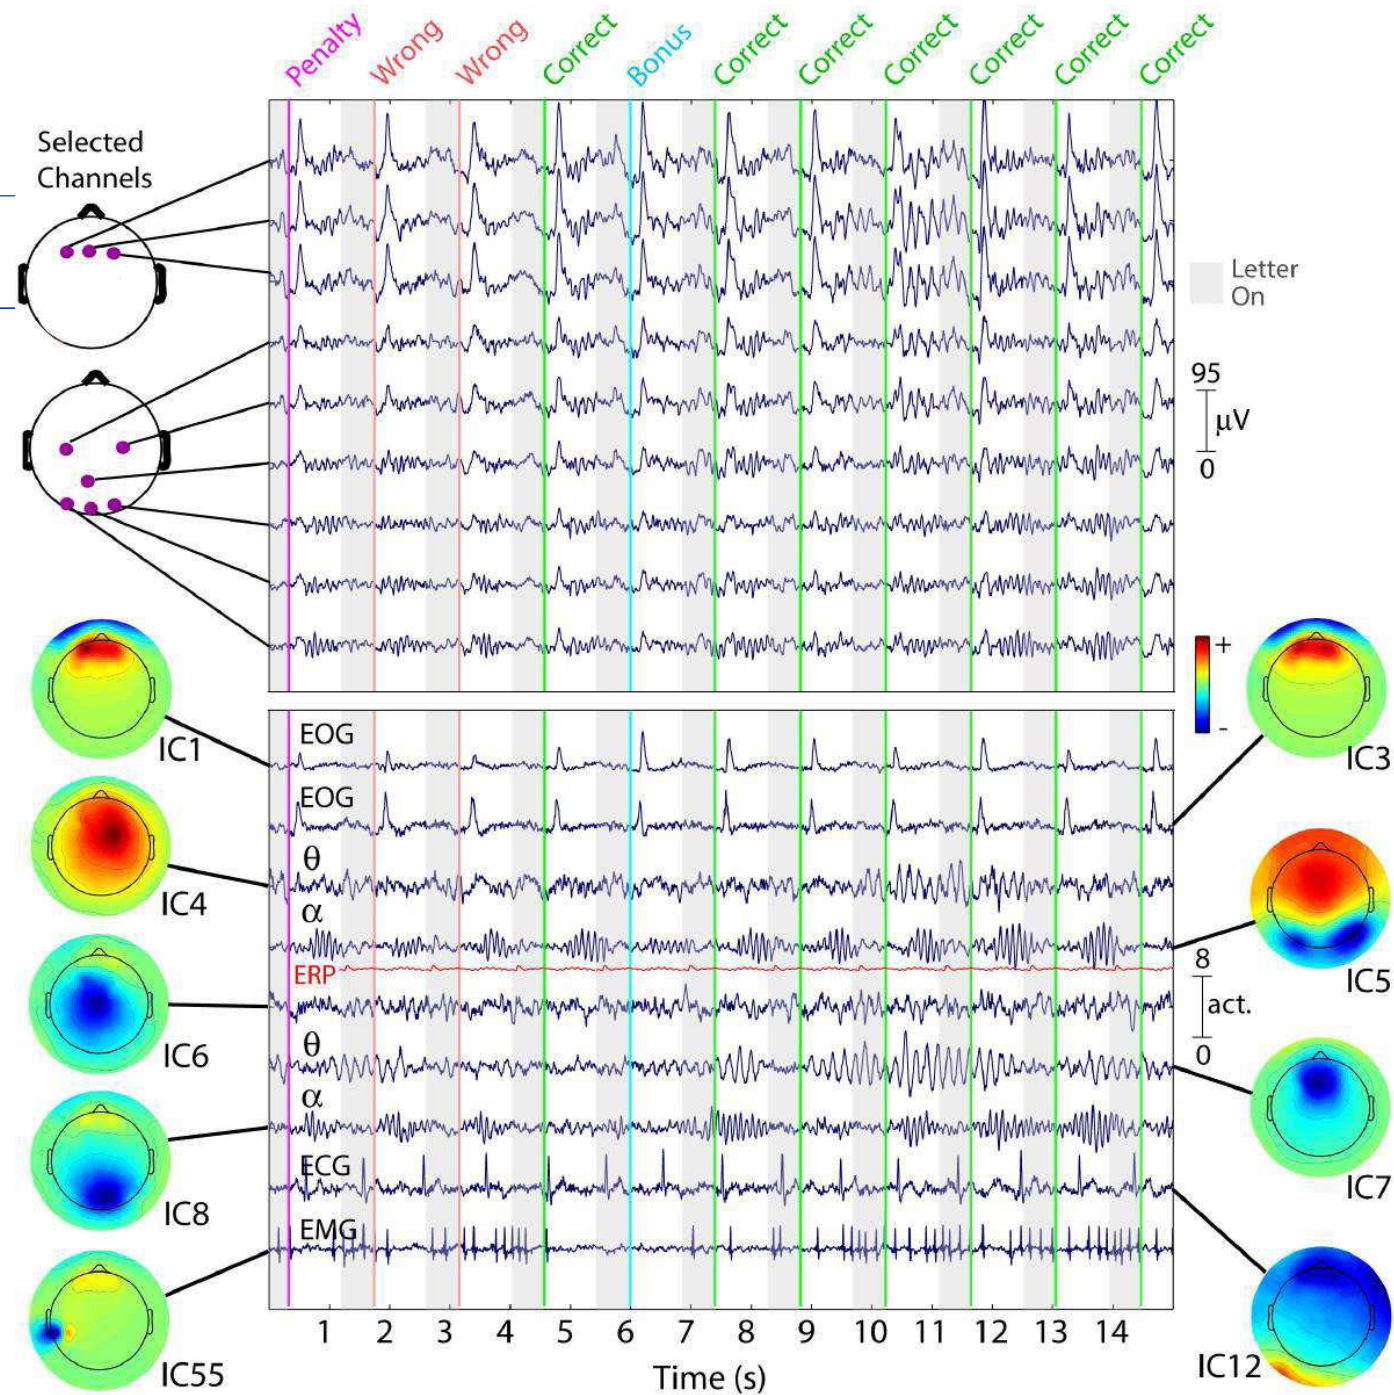
6. Independent Component Analysis (ICA)



- ICA can be used for handwritten digits



- Example:
- EEG feature extraction



Unsupervised Learning



1. Introduction
2. Association Rules & Cluster Analysis
3. Outlier Detection
4. Self-Organizing Maps
5. Principal Components, Curves and Surfaces
6. Non-negative Matrix Factorization
7. Independent Component Analysis
- 8. Multidimensional Scaling**
- 9. Nonlinear Dimension Reduction**
- 10. The Google PageRank Algorithm**

7. Multi-dimensional Scaling (MDS)



- Objectives: Data dimension reduction: different from SOMs, PCurves, MDS is to preserve the pairwise distance



- Data:

$$D = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}, \mathbf{x}_i \in R^p \Rightarrow S = \{\mathbf{s}_1, \mathbf{s}_2, \dots, \mathbf{s}_N\}, \mathbf{s}_i \in R^q (q < p)$$

- Distance or Dissimilarity $d_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|$
- Goal: find low dimension data S such that

$$S_D(z_1, z_2, \dots, z_N) = \sqrt{\sum_{i \neq j} \left(d_{ij} - \|z_i - z_j\| \right)^2}$$

7. Multi-dimensional Scaling (MDS)




14.9 Nonlinear Dimension Reduction




- Problem?

Classical MDS



Local MDS



Unsupervised Learning



1. Introduction
2. Association Rules & Cluster Analysis
3. Outlier Detection
4. Self-Organizing Maps
5. Principal Components, Curves and Surfaces
6. Non-negative Matrix Factorization
7. Independent Component Analysis
- 8. Multidimensional Scaling**
- 9. Nonlinear Dimension Reduction**
- 10. The Google PageRank Algorithm**

14.9 Nonlinear Dimension Reduction



- Local MDS is to **preserve the ordering of the points** along the curve.
- Three new approaches to nonlinear dimension reduction and manifold mapping.
 - Isometric feature mapping (ISOMAP) (Tenenbaum et al., 2000)
 - Local linear embedding (Roweis and Saul, 2000)
 - Local MDS (Chen and Buja, 2008)

Isometric feature mapping (ISOMAP)



- Isometric feature mapping (ISOMAP) constructs a graph to approximate the geodesic distance between points along the manifold.
- For each data point we find **its neighbors**—points within some small Euclidean distance of that point.
- We construct a graph with an edge between any two neighboring points.
- The geodesic distance between any two points is then approximated by **the shortest path** between points on the graph.

Local linear embedding



- Local linear embedding takes a very different approach, trying to preserve **the local affine structure** of the high-dimensional data.
- Each data point is approximated by a **linear combination of neighboring points**. Then a lower dimensional representation is constructed that best preserves these local approximations.

Local linear embedding



- For each data point x_i in p dimensions, we find its K -nearest neighbors $N(i)$ in Euclidean distance.
- We approximate each point by an affine mixture of the points in its neighborhood:


$$\min_{w_{ik}} \left\| x_i - \sum_{k \in N(i)} w_{ik} x_k \right\|^2$$

$$w_{ik} = 0, \quad k \notin N(i); \quad \sum_{k=1}^N w_{ik} = 1$$

- We must have $K < p$.

LLE


Keep w_{ij} unchanged



Original Space

Find a set of z^i minimizing

$$\sum_i \left\| z^i - \sum_j w_{ij} z^j \right\|_2$$



New (Low-dim) Space

Local linear embedding



- Given w_{ik} , we find points $y_i \in R^d$ in a space of dimension $d < p$ to minimize

$$\min_{y_i} \sum_{i=1}^N \left\| y_i - \sum_{k=1} w_{ik} y_k \right\|^2$$

- In compact form:

$$\min_Y \text{tr} \left[(Y - WY)^T (Y - WY) \right] = \text{tr} \left[Y^T (I - W)^T (I - W) Y \right]$$

where $W \in R^{N \times N}$, $Y \in R^{N \times d}$. The solution is the trailing eigenvectors of

$$M = (I - W)^T (I - W)$$

- Since 1 is a trivial eigenvector with eigenvalue 0, we discard it and keep the next d . This has the side effect that $1^T Y = 0$

- Define N to be the symmetric set of nearby pairs of points; specifically a pair (i, i') is in N if point i is among the K -nearest neighbors of i' , or vice-versa.


$$S(\{z_i\}_{i=1}^N) = \sum_{i, i' \in N} \left\| d_{ii'} - \|z_i - z_{i'}\| \right\|^2 \\ + \sum_{i, i' \notin N} w \left\| D - \|z_i - z_{i'}\| \right\|^2$$

$$w \sim 1/D$$


- Take $\tau=2wD$, for example: $\tau=0.01$

$$S(\{z_i\}_{i=1}^N) = \sum_{i,i' \in N} \left\| d_{ii'} - \|z_i - z_{i'}\| \right\|^2 - \tau \sum_{i,i' \notin N} \|z_i - z_{i'}\|^2$$

Classical MDS




Local MDS



- LLE example

20X28 image

- Similar results by LMDS



Stochastic Neighbor Embedding



- For many real-world dataset with non-linear inherent structures(e.g. MNIST), both linear methods like PCA and classical manifold learning algorithms like *Isomap* and *LLE* fail.
- Why? How to solve it?

3 6 3 / 7 9 6 6 4 1
6 7 5 7 8 6 8 4 8 5
2 1 7 9 7 / 2 1 1 5
4 8 1 9 0 / 8 8 9 4
3 6 1 3 4 / 5 4 0
7 5 9 2 6 5 8 1 9 7
1 2 2 2 2 3 4 8 0
0 2 3 8 0 7 3 8 5 7
0 1 4 6 4 6 0 2 9 3
7 / 2 8 1 6 1 8 6 /



(a) Visualization by Isomap.



(b) Visualization by LLE.

Stochastic Neighbor Embedding



- Solution:
 - SNE (Hinton and Roweis, 2002)
 - t-SNE (van der Maaten and Hinton, 2008)
- Problem to be solved:
 - Measure of similarity inherited
 - How to define data distribution
- Technical Approach:
 - KL divergence between the two probability distributions at each point
- Algorithm:
 - Iteratively learn low-dimensional embedding by minimizing the cost function

Stochastic Neighbor Embedding



- For each data point $x_i \in R^p$, we find points $y_i \in R^d$ in a space of dimension $d < p$ to minimize **KL** divergence between two distributions.

- Similarity matrix at high dimension:

$$p_{j|i} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\sigma_i^2)}$$

- Similarity matrix at low dimension:

$$q_{j|i} = \frac{\exp(-\|y_i - y_j\|^2)}{\sum_{k \neq i} \exp(-\|y_i - y_k\|^2)}$$

- The cost function:

$$L = \sum_i KL(P_i | Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$


Stochastic Neighbor Embedding



- The cost function:
$$L = \sum_i KL(P_i | Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$
- It has gradient:

$$\frac{\partial L}{\partial y_i} = 2 \sum_j (y_i - y_j)(p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j})$$

- SNE suffers from the "crowding problem": The area of the 2D map that is available to accommodate moderately distant data points will not be large enough compared with the area available to accommodate nearby data points.



t-SNE: Student-t distribution SNE



- For each data point $x_i \in R^p$, we find points $y_i \in R^d$ in a space of dimension $d < p$ to minimize **KL** divergence between two distributions.
- Student-t distribution SNE:
 - A symmetrized version of the SNE cost function with simpler gradients.
 - A **Student-t distribution** to compute the similarity in the low-dimensional

- Symetry :

$$p_{ij} = \frac{p_{i|j} + p_{j|i}}{2}$$

$$p_{j|i} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\sigma_i^2)}$$

- A Student-t distribution for $y_i \in R^d$

$$q_{ij} = \frac{\left(1 + \|y_i - y_j\|^2\right)^{-1}}{\sum_{k \neq l} \left(1 + \|y_i - y_k\|^2\right)^{-1}}$$

$$p_{st}(z | p) = \frac{\Gamma\left(\frac{p+1}{2}\right) \left(1 + \|z\|^2 / p\right)^{-\frac{p+1}{2}}}{\sqrt{p\pi} \Gamma(p/2)}$$

Student t-Distribution




- Suppose that $\{x_k\}_{k=1}^N$ are N samples from Gaussian distribution

$$\bar{x} = \frac{1}{N} \sum_{k=1}^N x_k; \quad \bar{S}^2 = \frac{1}{N-1} \sum_{k=1}^N (x_k - \bar{x})(x_k - \bar{x})^T$$

- $\frac{\bar{x} - \mu}{\bar{S} / \sqrt{N}} = \frac{\bar{x} - \mu}{\sigma / \sqrt{N}} \frac{1}{\bar{S} / \sigma}$ follows the student t-distribution with N-1 dfs, i.e. $p=N-1$

$$p_{st}(z | p) = \frac{\Gamma\left(\frac{p+1}{2}\right) \left(1 + \|z\|^2 / p\right)^{-\frac{p+1}{2}}}{\sqrt{p\pi} \Gamma(p/2)}$$



t-SNE: Student-t distribution SNE




- Cost function:

$$L = \sum_i KL(P_i | Q_i) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

- The gradient
$$\frac{\partial L}{\partial y_i} = 4 \sum_{j \neq i} (p_{ij} - q_{ij}) \left(1 + \|y_i - y_j\|^2 \right)^{-1} (y_i - y_j)$$

- The heavy tails of the normalized Student-t kernel allow dissimilar input objects $x_i, x_j \in R^p$ to be modeled by low-dimensional counterparts $y_i, y_j \in R^d$ that are too far apart because q_{ij} is not very small for two embedded points that are far apart.

Comparisons



(a) Visualization by t-SNE.



(a) Visualization by Isomap.



(b) Visualization by Sammon mapping.




(b) Visualization by LLE.

Unsupervised Learning



1. Introduction
2. Association Rules & Cluster Analysis
3. Association Rules & Cluster Analysis
4. Self-Organizing Maps
5. Principal Components, Curves and Surfaces
6. Non-negative Matrix Factorization
7. Independent Component Analysis
8. Multidimensional Scaling
9. Nonlinear Dimension Reduction
- 10. The Google PageRank Algorithm**

The Google PageRank Algorithm



- We suppose that we have **N** web pages and wish to rank them in terms of importance.
- The PageRank algorithm considers a webpage to be **important if many other webpages point to it**.
- The linking webpages that point to a given page are not treated equally: the algorithm also takes into account both the **importance** (PageRank) of the linking pages and the **number of outgoing links** that they have.

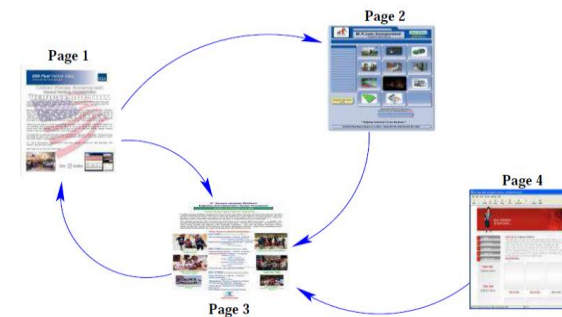


FIGURE 14.46. PageRank algorithm: example of a small network

PageRank Formulation



- Let $L_{ij} = 1$ if page j points to page i , and zero otherwise.
- Let $c_j = \sum_{i=1} L_{ij}$ equals the number of pages pointed to by page j (number of outlinks).

$$L_{12} = 1 \text{ or } L_{21} = 1 ?$$

$$c_1 = L_{11} + L_{21} + L_{31} + L_{41} = 2$$

$$c_3 = ?$$

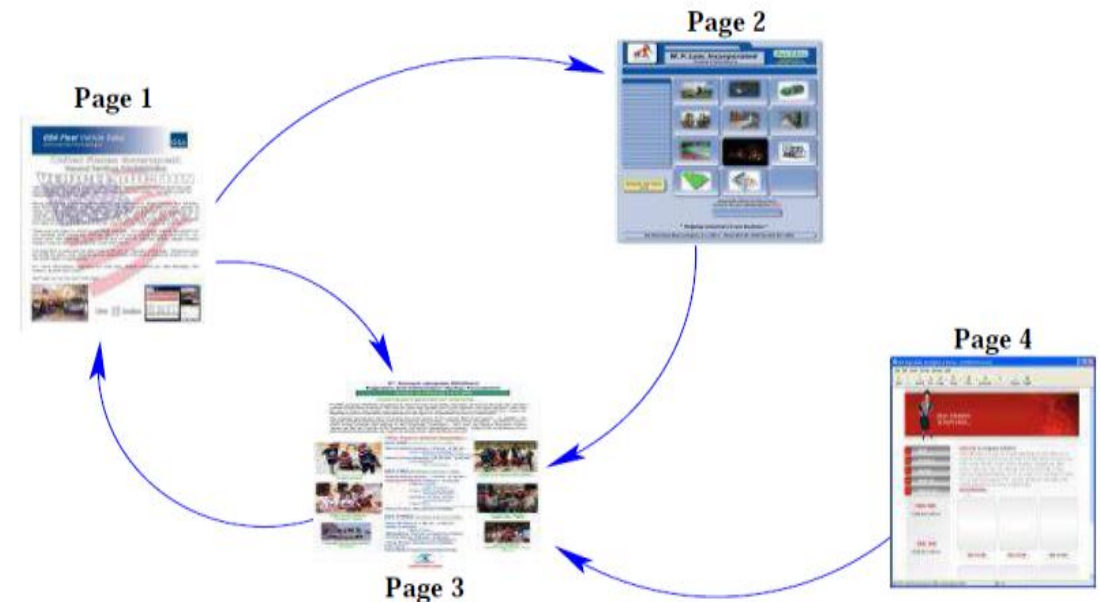


FIGURE 14.46. PageRank algorithm: example of a small network

PageRank Formulation



- **Assumption:** The importance of page i is the sum of the importance of other pages that point to that page.
- The **Google PageRanks** p_i are defined by the recursive relationship

$$p_i = (1 - d) + d \sum_{j=1}^N \frac{L_{ij}}{c_j} p_j,$$

d is a positive constant (say $d=0.85$). In matrix form

$$\mathbf{p} = (1 - d)\mathbf{e} + d\mathbf{L}\mathbf{D}_c^{-1}\mathbf{p}$$

\mathbf{e} is a vector of N ones; $\mathbf{D}_c = \text{diag}(c)$

Page Rank Algorithm



$$\mathbf{p} = (1 - d)\mathbf{e} + d\mathbf{LD}_c^{-1}\mathbf{p}$$

- Introducing the normalization $\mathbf{e}^T \mathbf{p} = N$ (i.e., the average PageRank is 1), rewrite the above equation

$$\mathbf{p} = \left[(1 - d)\mathbf{e}\mathbf{e}^T / N + d\mathbf{LD}_c^{-1} \right] \mathbf{p} = \mathbf{A}\mathbf{p}$$

- Exploiting a connection with Markov chains (see below), it can be shown that the matrix \mathbf{A} has a real eigenvalue equal to one, and one is its largest eigenvalue.

PageRank Algorithm



- The page rank solution

$$\mathbf{p} = \left[(1-d)\mathbf{e}\mathbf{e}^T / N + d\mathbf{L}\mathbf{D}_c^{-1} \right] \mathbf{p} = \mathbf{A}\mathbf{p}$$

$$\mathbf{e}^T \mathbf{p} = N$$

- Find by the power method: initialize p_0

$$\mathbf{p}_k \leftarrow \mathbf{A}\mathbf{p}_{k-1}; \quad \mathbf{p}_k \leftarrow N \frac{\mathbf{p}_k}{\mathbf{e}^T \mathbf{p}_k}$$

PageRank Algorithm



- Viewing PageRank as a Markov chain makes clear why the matrix \mathbf{A} has a maximal real eigenvalue of 1.
- Since \mathbf{A} has positive entries with each column summing to one, Markov chain theory tells us that it has a unique eigenvector with eigenvalue one, corresponding to the stationary distribution of the chain (Bremaud, 1999).

Simple Example

$$L = \begin{pmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

The number of outlinks $c = (2, 1, 1, 1)$.

$$p = (1.49, 0.78, 1.58, 0.15)$$

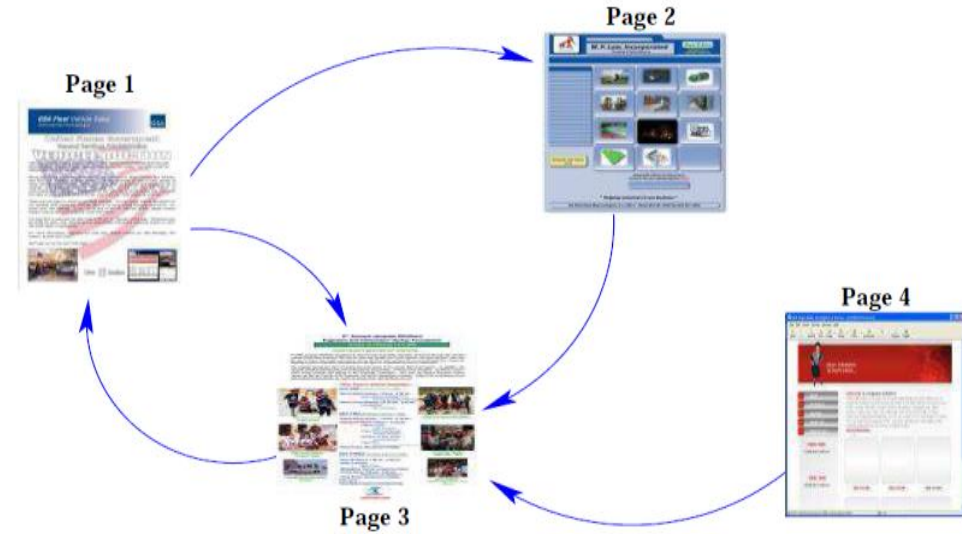


FIGURE 14.46. PageRank algorithm: example of a small network

1. Introduction
2. Association Rules & Cluster Analysis
3. Association Rules & Cluster Analysis
4. Self-Organizing Maps
5. Principal Components, Curves and Surfaces
6. Non-negative Matrix Factorization
7. Independent Component Analysis
8. Multidimensional Scaling
9. Nonlinear Dimension Reduction
- 10. The Google PageRank Algorithm**

The End of Talk